

HW1_JingQian_Q2

October 3, 2019

1 Load Spark

```
In [0]: # Install latest version of spark. If error, check the latest and replace "spark-2.4.4"
!apt-get install openjdk-8-jdk-headless -qq > /dev/null
!wget -q https://www-us.apache.org/dist/spark/spark-2.4.4/spark-2.4.4-bin-hadoop2.7.tgz
!tar xf spark-2.4.4-bin-hadoop2.7.tgz
!pip install -q findspark
import os
os.environ["JAVA_HOME"] = "/usr/lib/jvm/java-8-openjdk-amd64"
os.environ["SPARK_HOME"] = "/content/spark-2.4.4-bin-hadoop2.7"
import findspark
findspark.init()
```

```
In [0]: #The entry point to using Spark SQL is an object called SparkSession.
#It initiates a Spark Application which all the code for that Session will run on.
from pyspark.sql import SparkSession
spark = SparkSession.builder \
    .master("local[*]") \
    .appName("Learning_Spark") \
    .getOrCreate()
```

```
In [0]: import numpy as np
import pandas as pd

from google.colab import drive
drive.mount('/content/gdrive')
```

Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6b

Enter your authorization code:

ûûûûûûûûûûû

Mounted at /content/gdrive

2 Q2. Binary classification with Spark MLlib

2.1 (1). Data loading

Read the csv file into a Dataframe. You could set "inferschema" to true and rename the columns with the following information: "age", "workclass", "fnlwgt", "education", "education_num", "marital_status", "occupation", "relationship", "race", "sex", "capital_gain", "capital_loss", "hours_per_week", "native_country", "income".

```
In [0]: DATA_PATH = "/content/gdrive/My Drive/BigData/q1/adult_data.csv"
```

```
In [74]: data = spark.read.csv(DATA_PATH,inferSchema=True, header=False)
         print(data.count(),len(data.columns))
         data.show(5)
```

```
32561 15
```

	_c0	_c1	_c2	_c3	_c4	_c5	_c6
39	State-gov	77516.0	Bachelors	13.0	Never-married	Adm-clerical	Not-in-f
50	Self-emp-not-inc	83311.0	Bachelors	13.0	Married-civ-spouse	Exec-managerial	Hu
38	Private	215646.0	HS-grad	9.0	Divorced	Handlers-cleaners	Not-in-f
53	Private	234721.0	11th	7.0	Married-civ-spouse	Handlers-cleaners	Hu
28	Private	338409.0	Bachelors	13.0	Married-civ-spouse	Prof-specialty	

only showing top 5 rows

```
In [75]: col_names = ["age", "workclass", "fnlwgt", "education", "education_num", "marital_status",
                    "occupation", "relationship", "race", "sex", "capital_gain", "capital_loss",
                    "hours_per_week", "native_country", "income"]
         print(len(col_names))
         data = data.toDF(*col_names)
         data.show(5)
```

```
15
```

	age	workclass	fnlwgt	education	education_num	marital_status	occupation
39	State-gov	77516.0	Bachelors		13.0	Never-married	Adm-clerical
50	Self-emp-not-inc	83311.0	Bachelors		13.0	Married-civ-spouse	Exec-managerial
38	Private	215646.0	HS-grad		9.0	Divorced	Handlers-cleaners
53	Private	234721.0	11th		7.0	Married-civ-spouse	Handlers-cleaners
28	Private	338409.0	Bachelors		13.0	Married-civ-spouse	Prof-specialty

only showing top 5 rows

```
In [76]: data.printSchema()
```

```
root
 |-- age: integer (nullable = true)
 |-- workclass: string (nullable = true)
 |-- fnlwgt: double (nullable = true)
 |-- education: string (nullable = true)
 |-- education_num: double (nullable = true)
 |-- marital_status: string (nullable = true)
 |-- occupation: string (nullable = true)
 |-- relationship: string (nullable = true)
 |-- race: string (nullable = true)
 |-- sex: string (nullable = true)
 |-- capital_gain: double (nullable = true)
 |-- capital_loss: double (nullable = true)
 |-- hours_per_week: double (nullable = true)
 |-- native_country: string (nullable = true)
 |-- income: string (nullable = true)
```

2.2 (2). Data preprocessing

Convert the categorical variables into numeric variables with ML Pipelines and Feature Transformers. You will probably need OneHotEncoderEstimator, StringIndexer, and VectorAssembler. Split your data into training set and test set with ratio of 70% and 30% and set the seed to 100.

Reference: <https://towardsdatascience.com/machine-learning-with-pyspark-and-mllib-solving-a-binary-classification-problem-96396065d2aa>

```
In [77]: train, test = data.randomSplit([0.7, 0.3], seed = 100)
        print("Training Dataset Count: " + str(train.count()))
        print("Test Dataset Count: " + str(test.count()))
```

```
Training Dataset Count: 22838
```

```
Test Dataset Count: 9723
```

```
In [0]: from pyspark.ml.feature import OneHotEncoderEstimator, StringIndexer, VectorAssembler
        categoricalColumns = ['workclass', 'education', 'marital_status', 'occupation', 'relationship',
                              'race', 'sex', 'native_country']

        stages = []
        for categoricalCol in categoricalColumns:
            stringIndexer = StringIndexer(inputCol = categoricalCol, outputCol = categoricalCol)
            encoder = OneHotEncoderEstimator(inputCols=[stringIndexer.getOutputCol()], outputCols=[stringIndexer.getOutputCol()])
            stages += [stringIndexer, encoder]

        label_stringIdx = StringIndexer(inputCol = 'income', outputCol = 'label')
        stages += [label_stringIdx]

        numericCols = ['age', 'fnlwgt', 'education_num', 'capital_gain', 'capital_loss', 'hours_per_week']
        assemblerInputs = [c + "classVec" for c in categoricalColumns] + numericCols
```

```

assembler = VectorAssembler(inputCols=assemblerInputs, outputCol="features")
stages += [assembler]

```

```

In [79]: from pyspark.ml import Pipeline
         pipeline = Pipeline(stages = stages)
         pipelineModel = pipeline.fit(train)
         train = pipelineModel.transform(train)

         selectedCols = ['label', 'features'] + col_names
         train = train.select(selectedCols)
         train.printSchema()

```

```

root
|-- label: double (nullable = false)
|-- features: vector (nullable = true)
|-- age: integer (nullable = true)
|-- workclass: string (nullable = true)
|-- fnlwt: double (nullable = true)
|-- education: string (nullable = true)
|-- education_num: double (nullable = true)
|-- marital_status: string (nullable = true)
|-- occupation: string (nullable = true)
|-- relationship: string (nullable = true)
|-- race: string (nullable = true)
|-- sex: string (nullable = true)
|-- capital_gain: double (nullable = true)
|-- capital_loss: double (nullable = true)
|-- hours_per_week: double (nullable = true)
|-- native_country: string (nullable = true)
|-- income: string (nullable = true)

```

2.3 (3). Modelling

Train a logistic regression model with train set. Learn more about models provide in Spark MLlib [here](#) . After training, plot ROC curve and Precision-Recall curve of your training process.

```

In [0]: from pyspark.ml.classification import LogisticRegression
         lr = LogisticRegression(featuresCol = 'features', labelCol = 'label', maxIter=10)
         lrModel = lr.fit(train)

```


```

In [81]: import matplotlib.pyplot as plt
         trainingSummary = lrModel.summary
         roc = trainingSummary.roc.toPandas()
         plt.plot(roc['FPR'],roc['TPR'])
         plt.ylabel('False Positive Rate')
         plt.xlabel('True Positive Rate')
         plt.title('ROC Curve')


```

```
plt.show()
print('Training set areaUnderROC: ' + str(trainingSummary.areaUnderROC))

pr = trainingSummary.pr.toPandas()
plt.plot(pr['recall'],pr['precision'])
plt.ylabel('Precision')
plt.xlabel('Recall')
plt.show()
```



Training set areaUnderROC: 0.9056654937412549



2.4 (4). Evaluation

Apply your trained model on the test set. Provide the value of area under ROC, accuracy, and confusion matrix. You should expect the accuracy to be around 85%.

```
In [82]: test = pipelineModel.transform(test)
        test = test.select(selectedCols)

        predictions = lrModel.transform(test)

        from pyspark.ml.evaluation import BinaryClassificationEvaluator
        evaluator = BinaryClassificationEvaluator()
        print('Test Area Under ROC', evaluator.evaluate(predictions))
```

Test Area Under ROC 0.9027382028865563

```
In [84]: from pyspark.ml.evaluation import MulticlassClassificationEvaluator

        # Select (prediction, true label) and compute test error
        evaluator = MulticlassClassificationEvaluator(
            labelCol="label", predictionCol="prediction", metricName="accuracy")
        accuracy = evaluator.evaluate(predictions)
        print("Test accuracy: ", accuracy)
```

Test accuracy: 0.8484006993726216

```
In [87]: from sklearn.metrics import confusion_matrix
          y_true = test.select('label').collect()
          y_pred = predictions.select('prediction').collect()
          cnf_matrix = confusion_matrix(y_true, y_pred)
          cnf_matrix
```

```
Out[87]: array([[6860,  530],
                [ 944, 1389]])
```